


Cyberbullying Detection on Social Media

Anurag Rajendra Watane


11013614

What is Cyberbullying?

- The term “Cyberbullying” means, use of Information Technology to harm or harass other people in a deliberate, repeated, and hostile manner
- Mean text messages, spreading rumours, posting messages, and sharing embarrassing photos and videos on social networking sites are all examples of cyberbullying




Twitter Dataset



Combined Data

Total 13485 rows

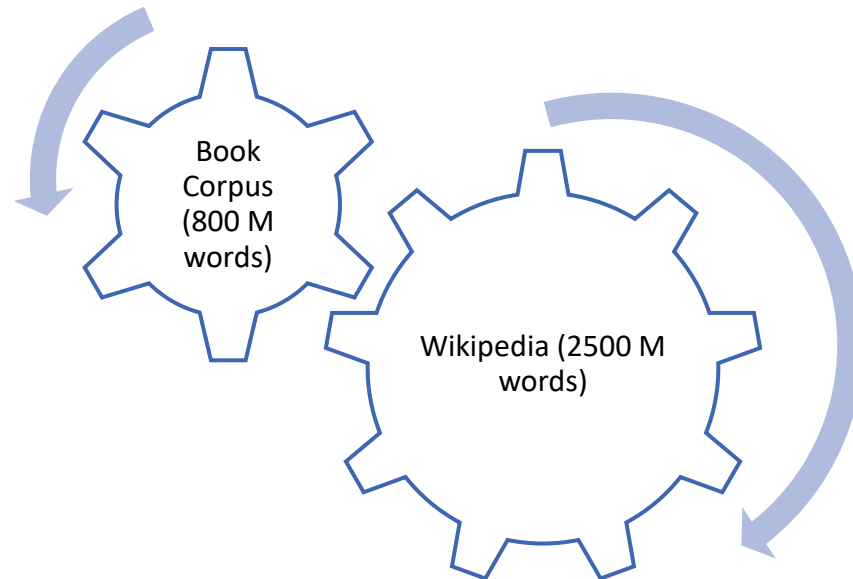


[HASOC \(hasocfire.github.io\)](https://hasocfire.github.io)

[UMICH SI650 - Sentiment Classification](#) | [Kaggle](#)

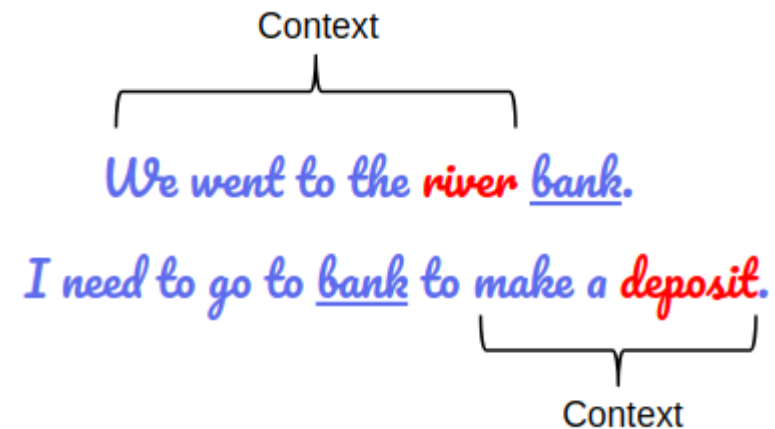
What is BERT?

- Bidirectional Encoder Representations from Transformers
- It is based on the Transformer architecture
- BERT is pre-trained on a large corpus of unlabelled text



Why Bidirectional?

- It uses multilayer bidirectional transformer encoders for language representations
- BERT is a “deeply bidirectional” model. Bidirectional means that BERT learns information from both the left and the right side of a token’s context during the training phase

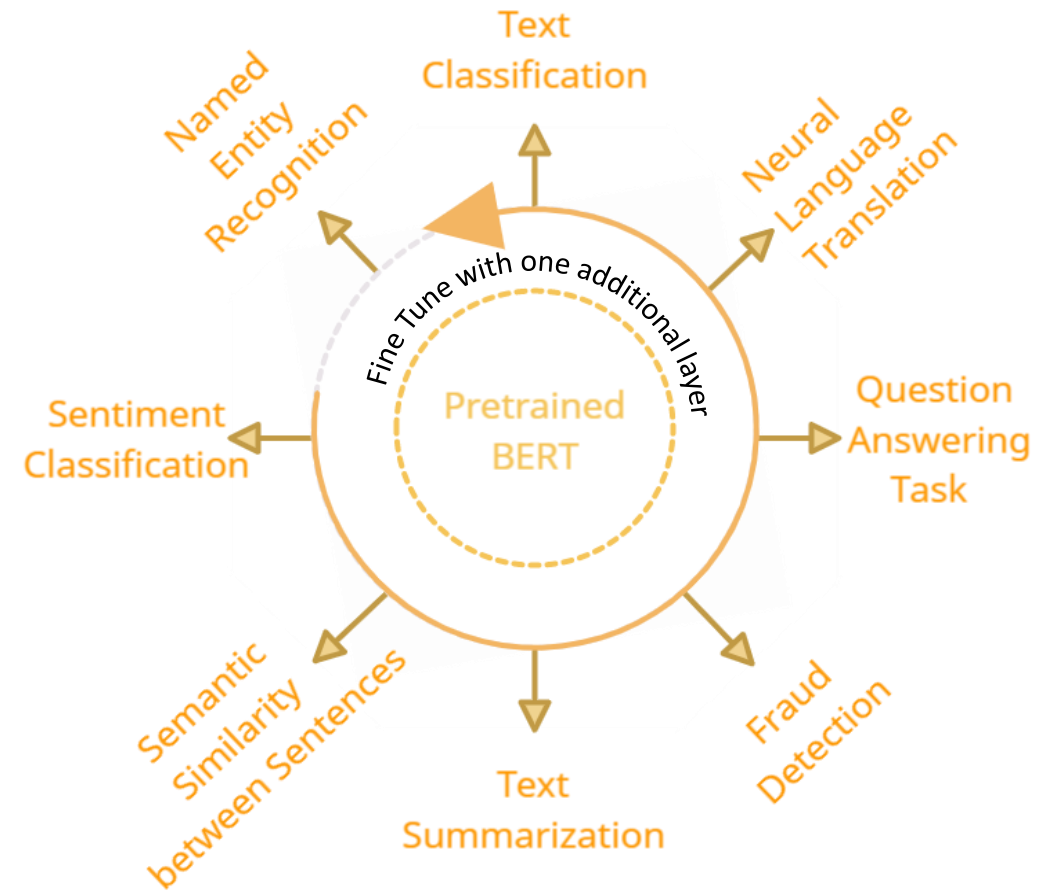


Why pretrained model?

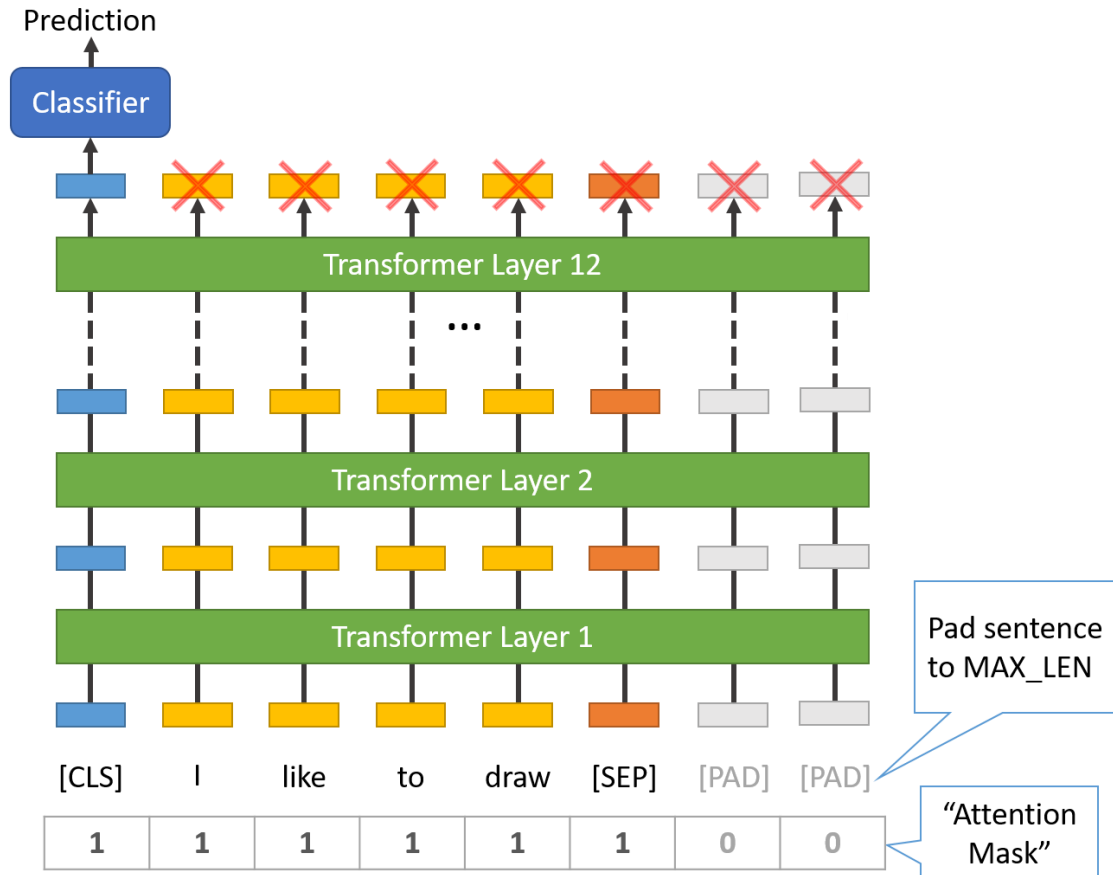
- The pre-trained BERT model weights already encode a lot of information about our language
- Because of the pre-trained weights this method allows us to fine-tune our task on a much smaller dataset
- It takes much less time to fine-tune the model

Advantages of pre-trained model

- BERT uses the concept of fine-tuning, and the final model for any task is almost the same as BERT.

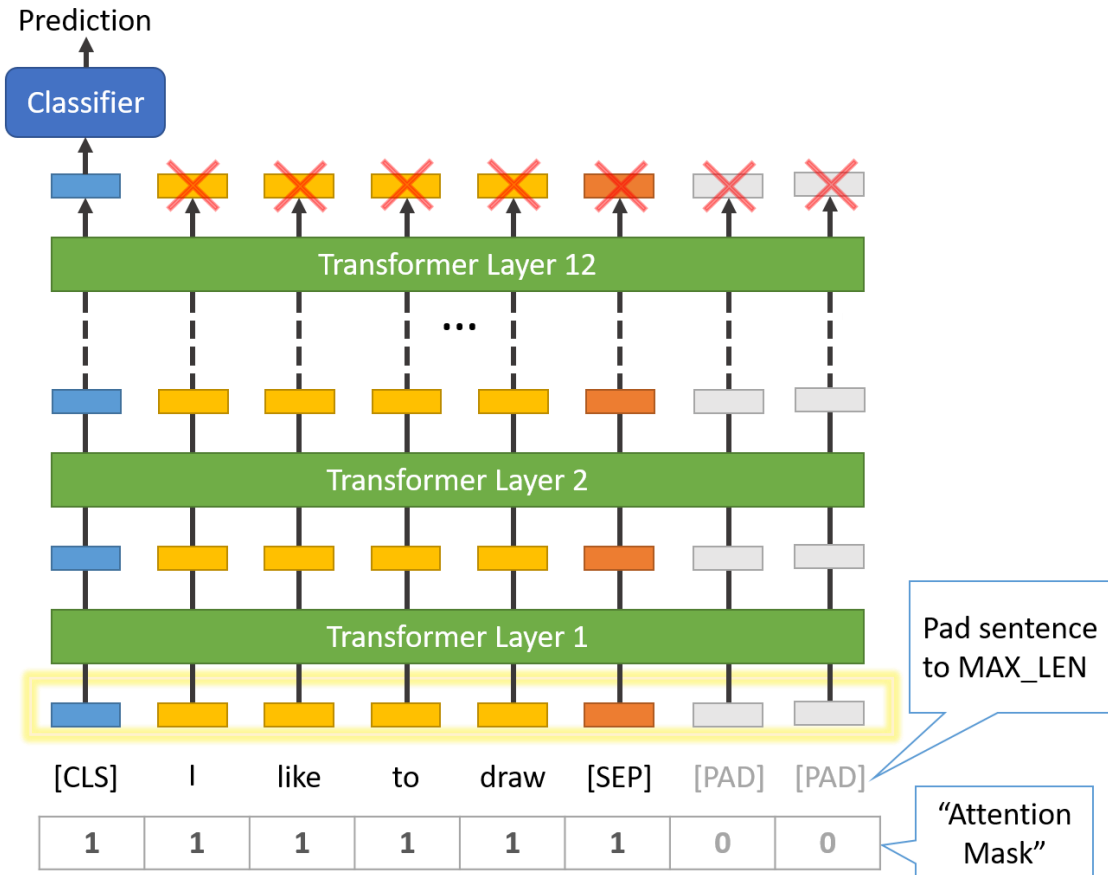


Architecture Used



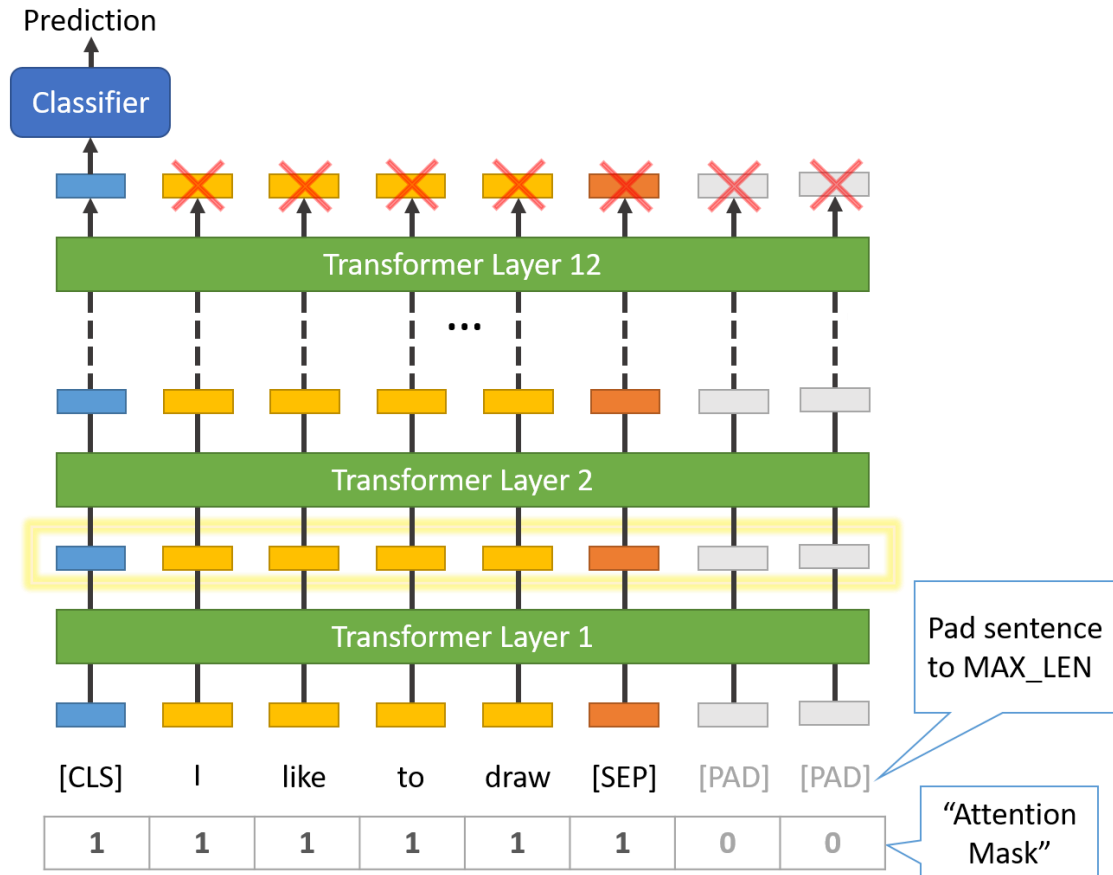
- 12-layer, 768-hidden, 12-heads, 110M parameters
- I have used pre-trained weights of 'bert-base-uncased' to fine-tune the model on Twitter dataset
- This model is uncased: it does not make a difference between **e**nglish and **E**nglish

Architecture Used



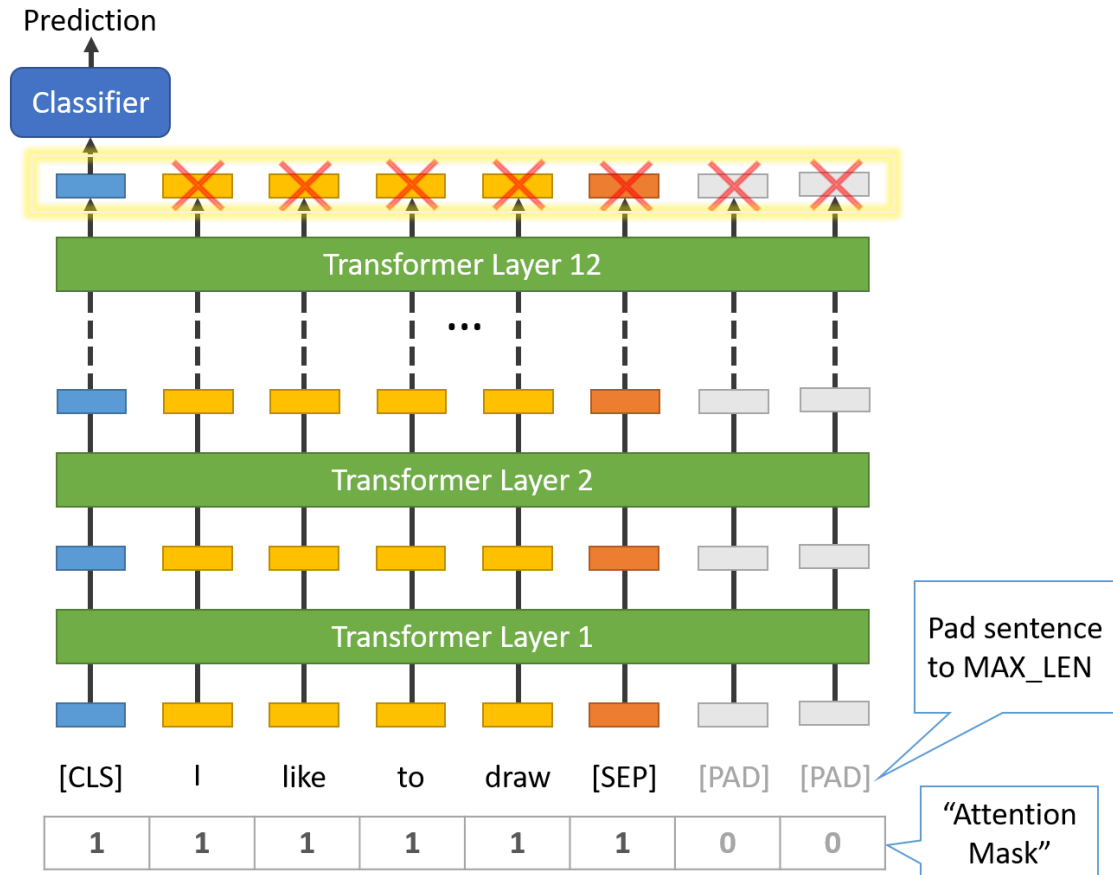
- 12-layer, 768-hidden, 12-heads, 110M parameters
- Pre-trained weights of 'bert-base-uncased' are used to fine-tune the model on Twitter dataset
- This model is uncased: it does not make a difference between **e**nglish and **E**nglish

Architecture Used



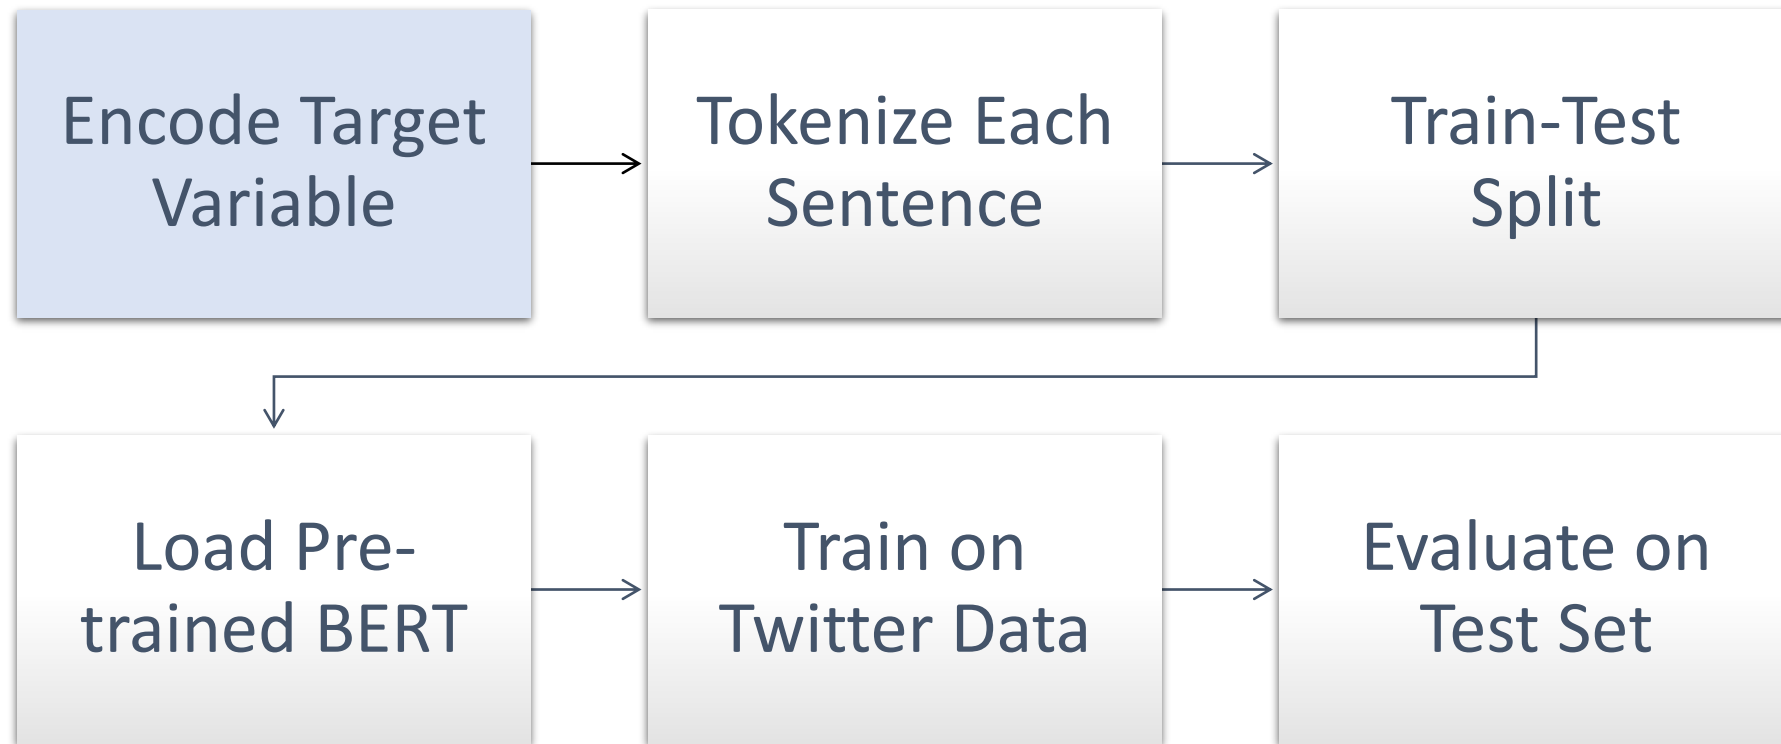
- 12-layer, 768-hidden, 12-heads, 110M parameters
- I have used pre-trained weights of 'bert-base-uncased' to fine-tune the model on Twitter dataset
- This model is uncased: it does not make a difference between **e**nglish and **E**nglish

Architecture Used

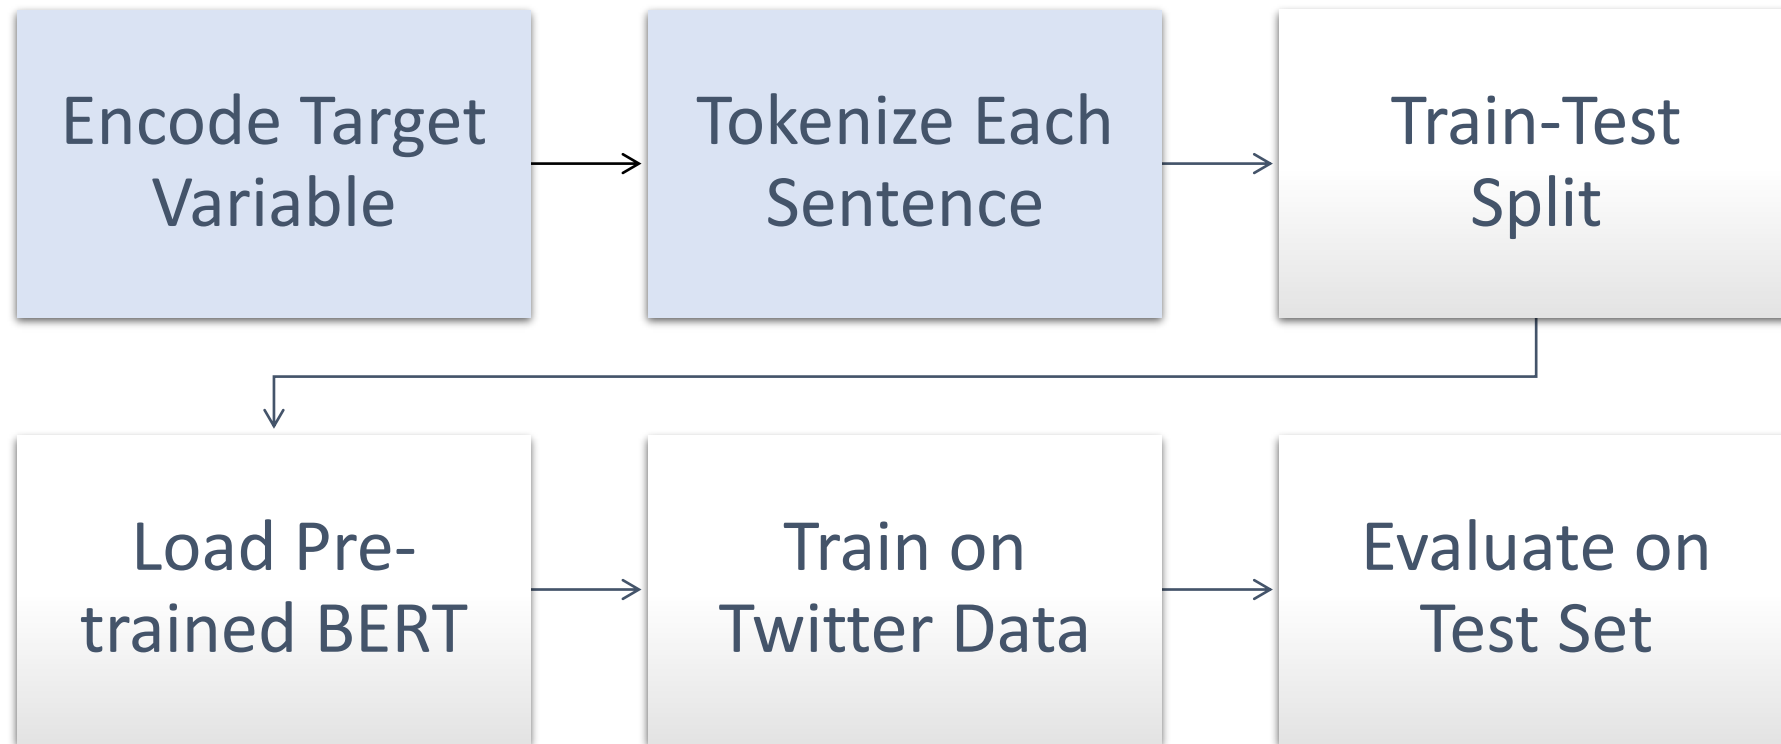


- 12-layer, 768-hidden, 12-heads, 110M parameters
- I have used pre-trained weights of 'bert-base-uncased' to fine-tune the model on Twitter dataset
- This model is uncased: it does not make a difference between **e**nglish and **E**nglish

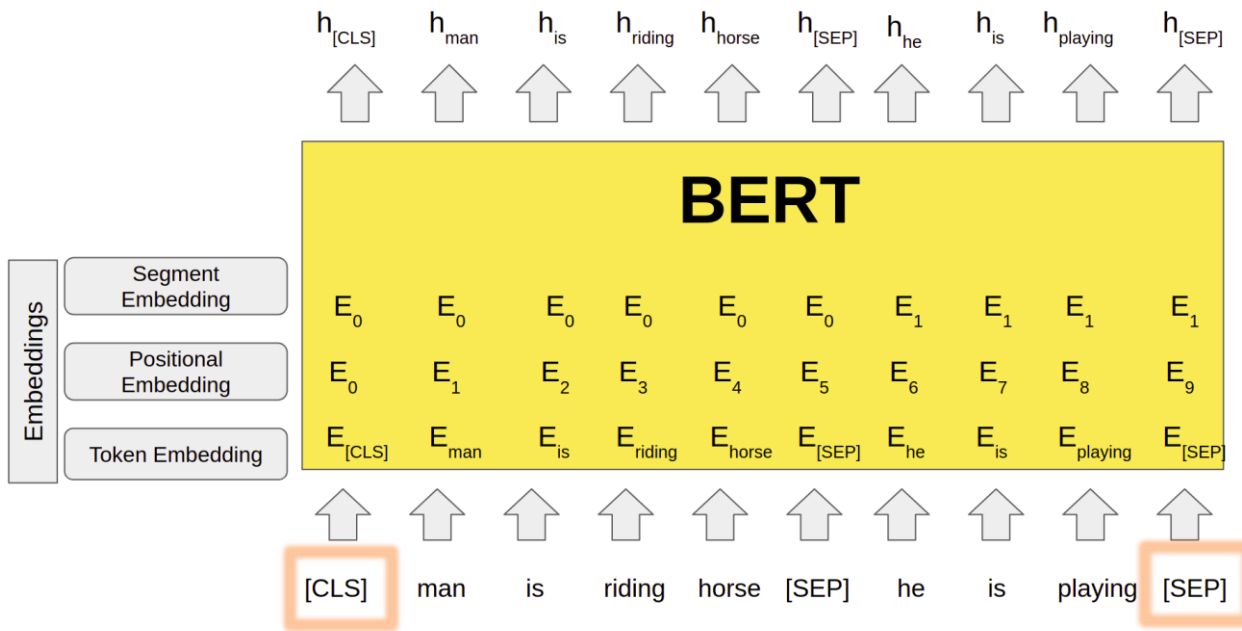
Steps for fine-tuning BERT



Steps for fine-tuning BERT

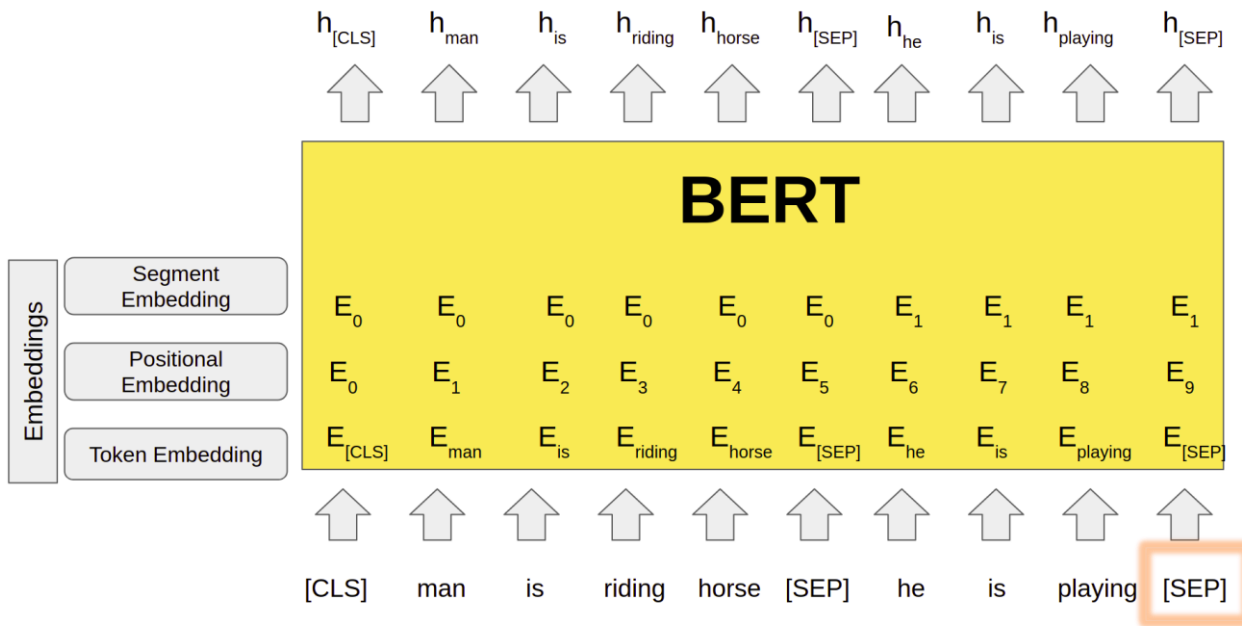


Tokenization for BERT



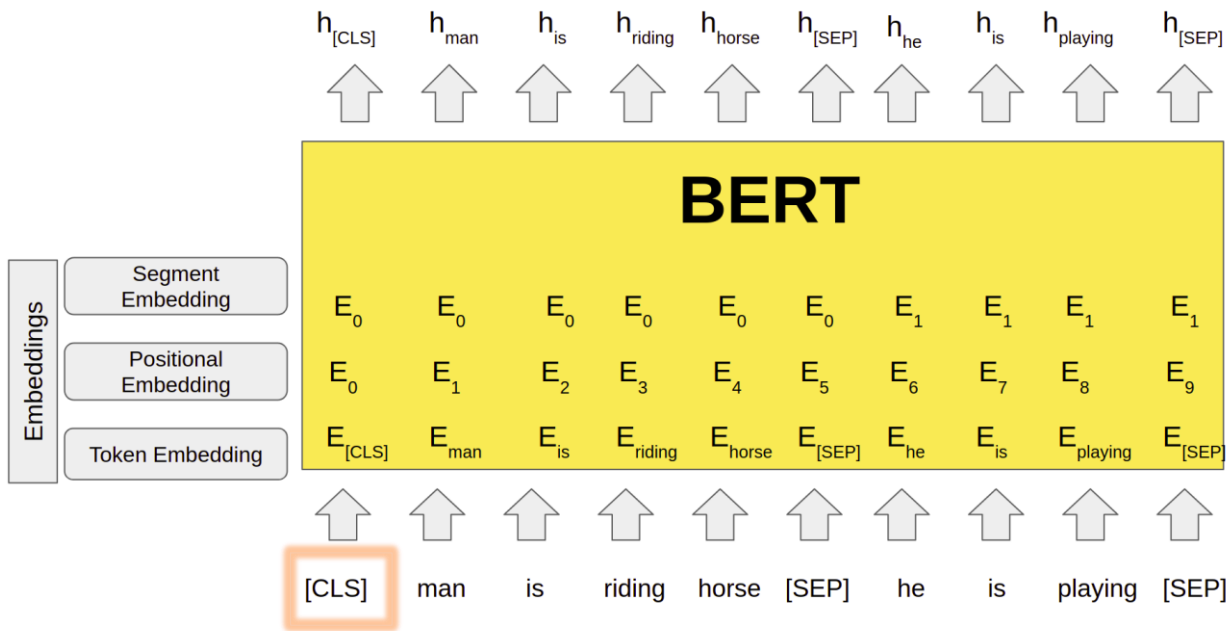
- BERT uses special tokens [CLS] and [SEP] to understand input properly
- [SEP] token has to be inserted at the end of every sentence
- The first token of every input sequence is the special classification token – [CLS]
- The last hidden state of BERT, corresponding to this token ($h_{[CLS]}$) is used in classification tasks as an aggregate of the entire sequence representation

Tokenization for BERT



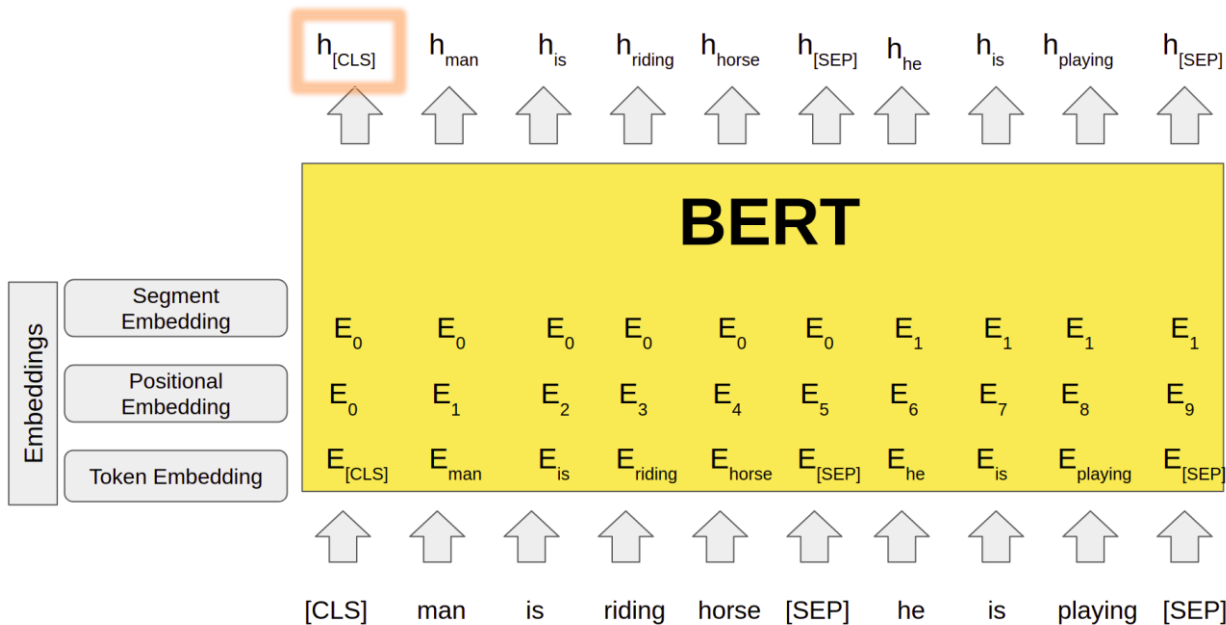
- BERT uses special tokens [CLS] and [SEP] to understand input properly
- [SEP] token has to be inserted at the end of every sentence
- The first token of every input sequence is the special classification token – [CLS]
- The last hidden state of BERT, corresponding to this token ($h[CLS]$) is used in classification tasks as an aggregate of the entire sequence representation

Tokenization for BERT




- BERT uses special tokens [CLS] and [SEP] to understand input properly
- [SEP] token has to be inserted at the end of every sentence
- The first token of every input sequence is the special classification token – [CLS]
- The last hidden state of BERT, corresponding to this token ($h[CLS]$) is used in classification tasks as an aggregate of the entire sequence representation

Tokenization for BERT



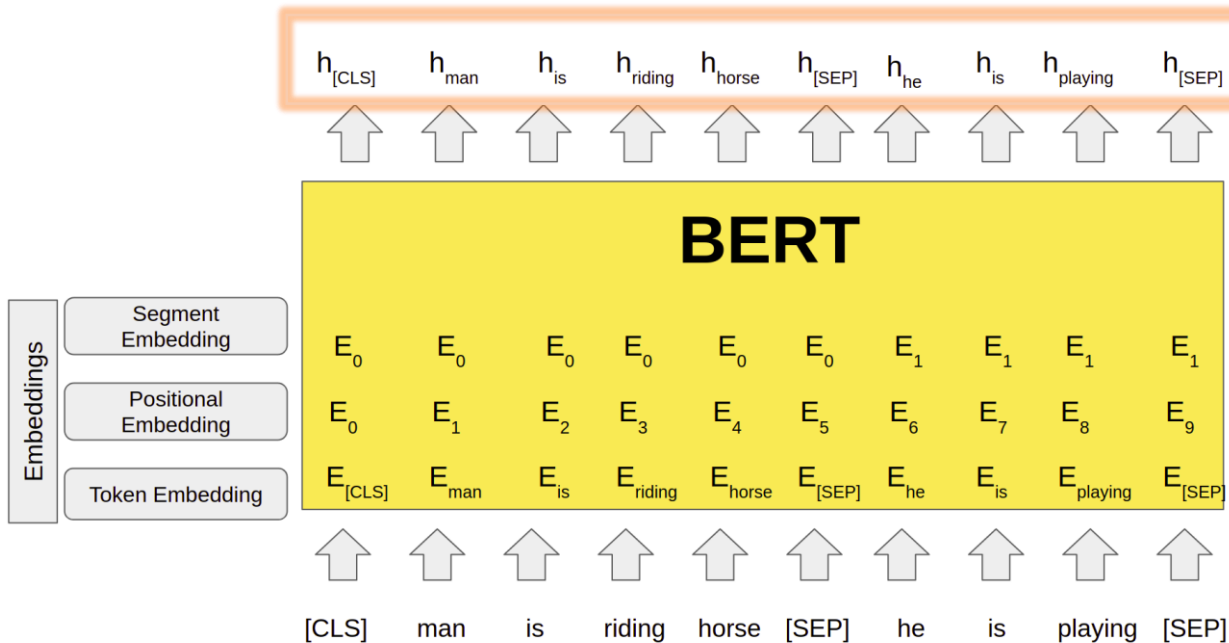
- BERT uses special tokens [CLS] and [SEP] to understand input properly
- [SEP] token has to be inserted at the end of every sentence
- The first token of every input sequence is the special classification token – [CLS]
- The last hidden state of BERT, corresponding to this token ($h[CLS]$) is used in classification tasks as an aggregate of the entire sequence representation

Tokenization for BERT



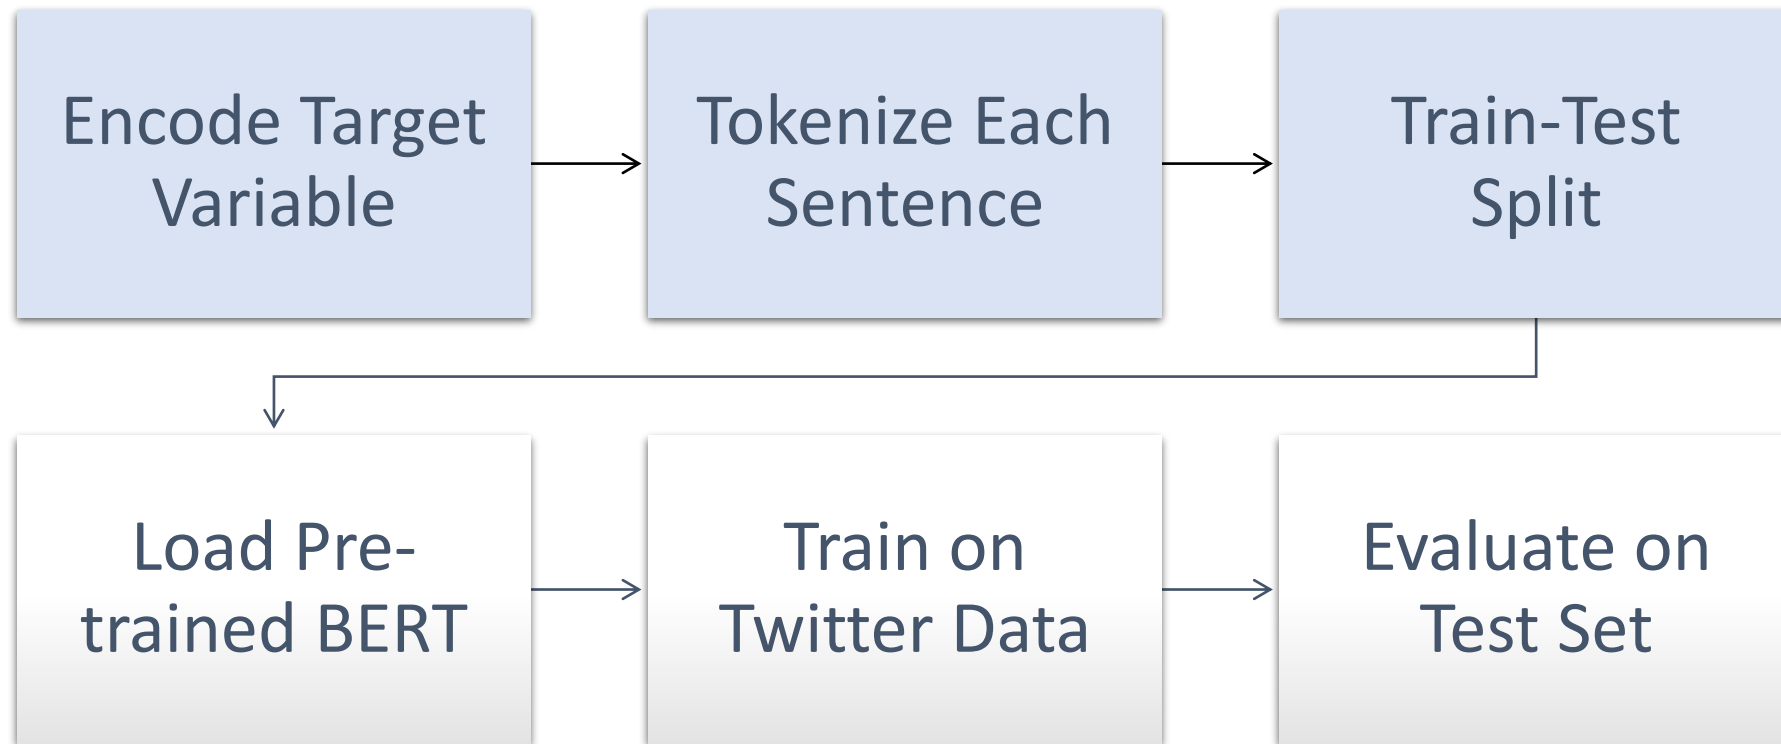
- BERT uses special tokens [CLS] and [SEP] to understand input properly
- [SEP] token has to be inserted at the end of every sentence
- The first token of every input sequence is the special classification token – [CLS]
- The last hidden state of BERT, corresponding to this token ($h_{[CLS]}$) is used in classification tasks as an aggregate of the entire sequence representation

Tokenization for BERT

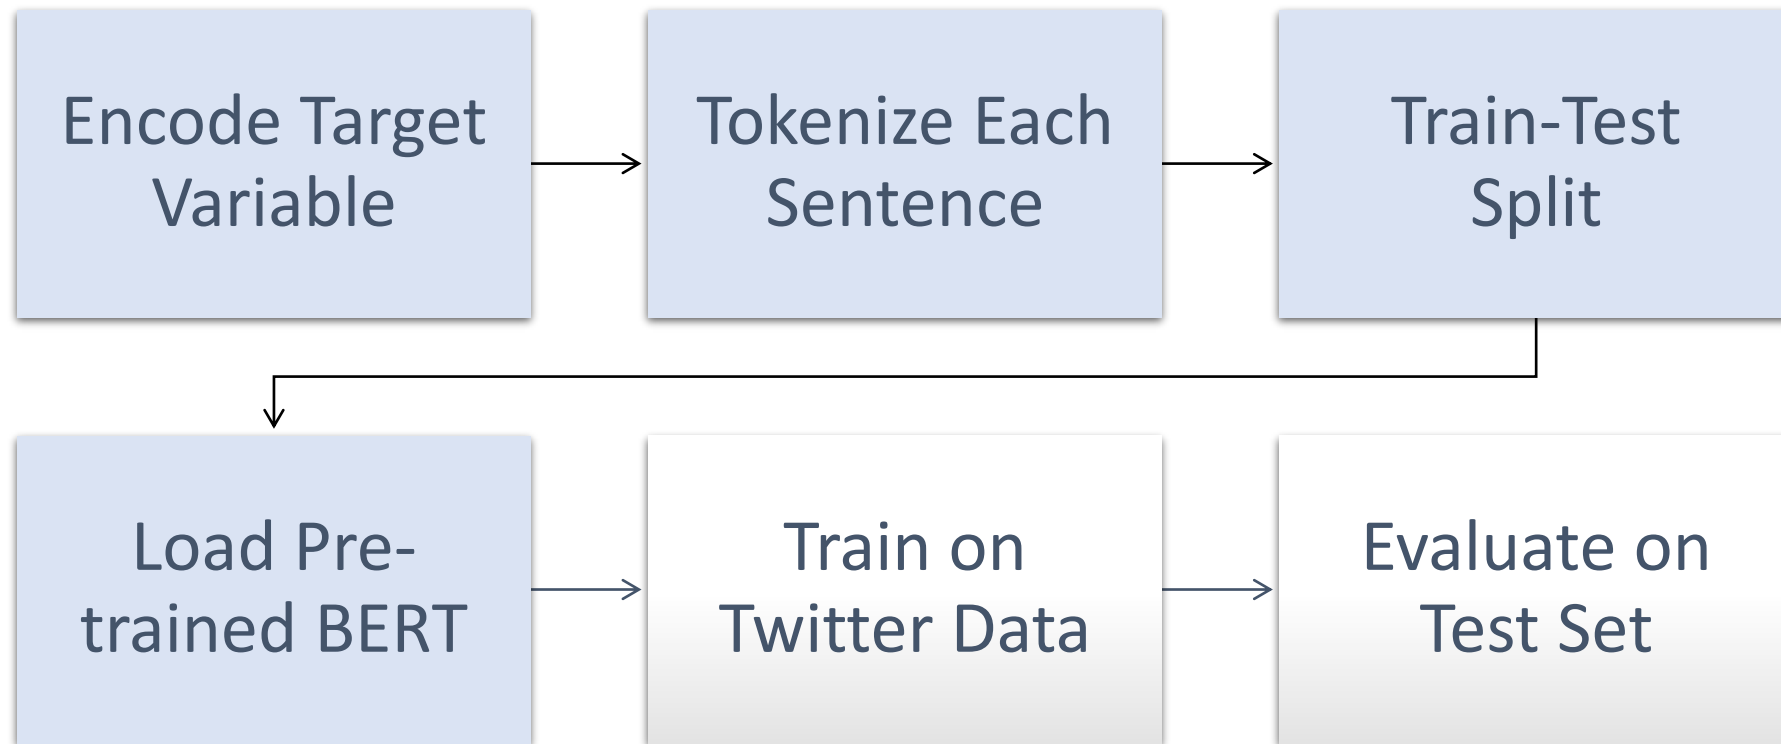


- BERT uses special tokens [CLS] and [SEP] to understand input properly
- [SEP] token has to be inserted at the end of every sentence
- The first token of every input sequence is the special classification token – [CLS]
- The last hidden state of BERT, corresponding to this token ($h_{[CLS]}$) is used in classification tasks as an aggregate of the entire sequence representation

Steps for fine-tuning BERT



Steps for fine-tuning BERT

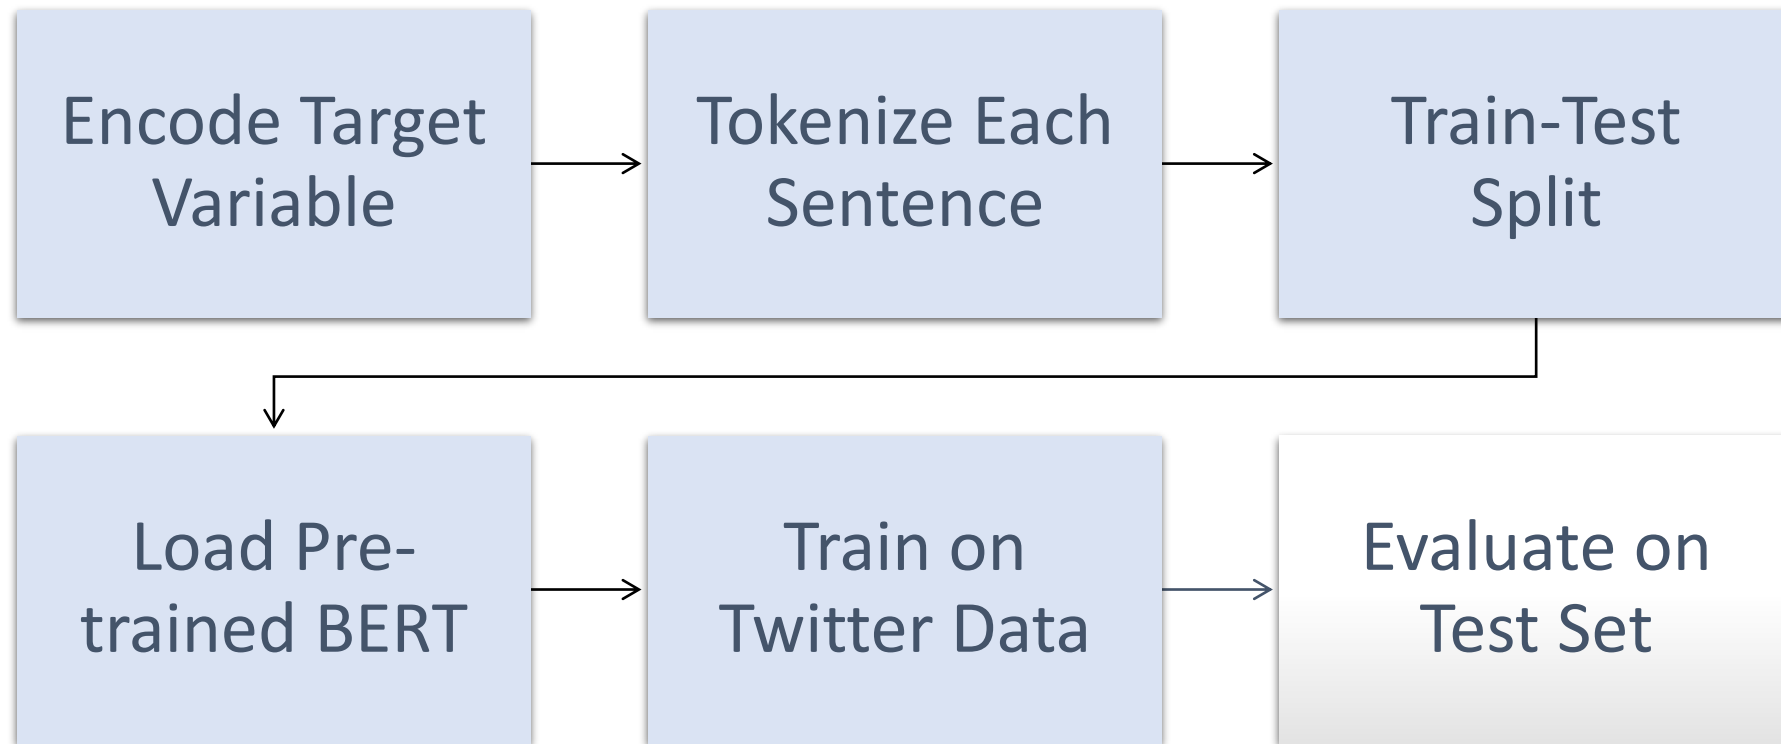


Load pretrained BERT

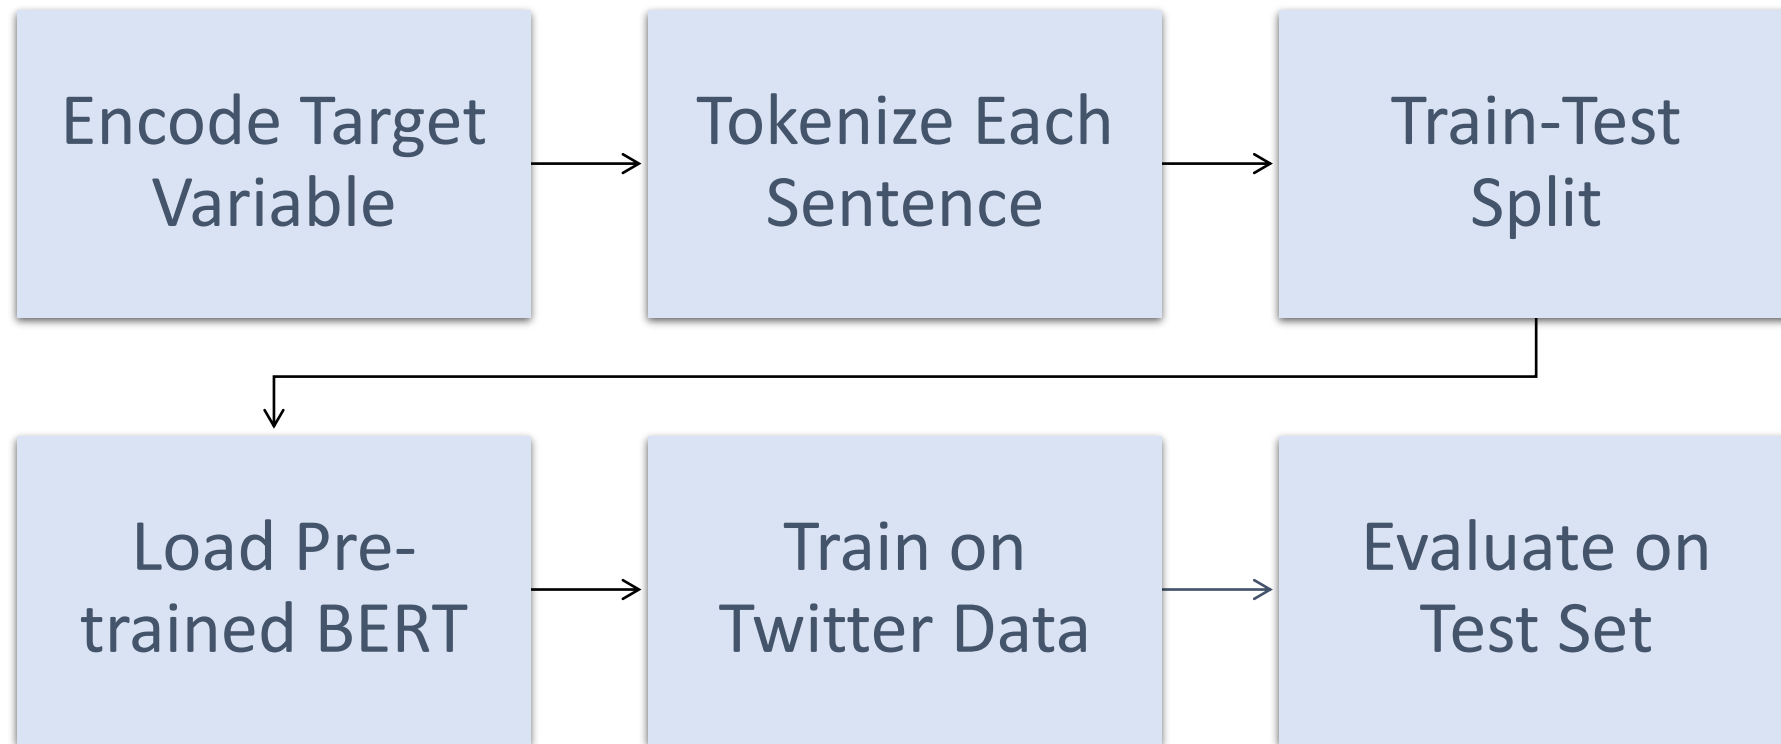
```
from transformers import BertForSequenceClassification
model = BertForSequenceClassification.from_pretrained(
    "bert-base-uncased",          # Use the 12-layer BERT model, with an uncased vocab
    num_labels = 2,              # The number of output labels--2 for binary classification
    output_attentions = False,    # Whether the model returns attentions weights
    output_hidden_states = False, # Whether the model returns all hidden-states
)
```

- BertForSequenceClassification is used to load the pre-trained BERT model from 'transformers' library developed by company Hugging Face
- The BertForSequenceClassification is used to fine-tune the BERT model
- BertForSequenceClassification is the normal BERT model with an added single linear layer on top for classification

Steps for fine-tuning BERT



Steps for fine-tuning BERT



F1 Score: BERT

```
accurate = 0
for (i,j) in zip(flat_predictions, flat_true_labels):
    if i==j:
        accurate += 1
accurate/len(flat_predictions)
```

```
0.8894348894348895
```

```
from sklearn.metrics import f1_score
f1_score(flat_true_labels, flat_predictions, average='macro')
```

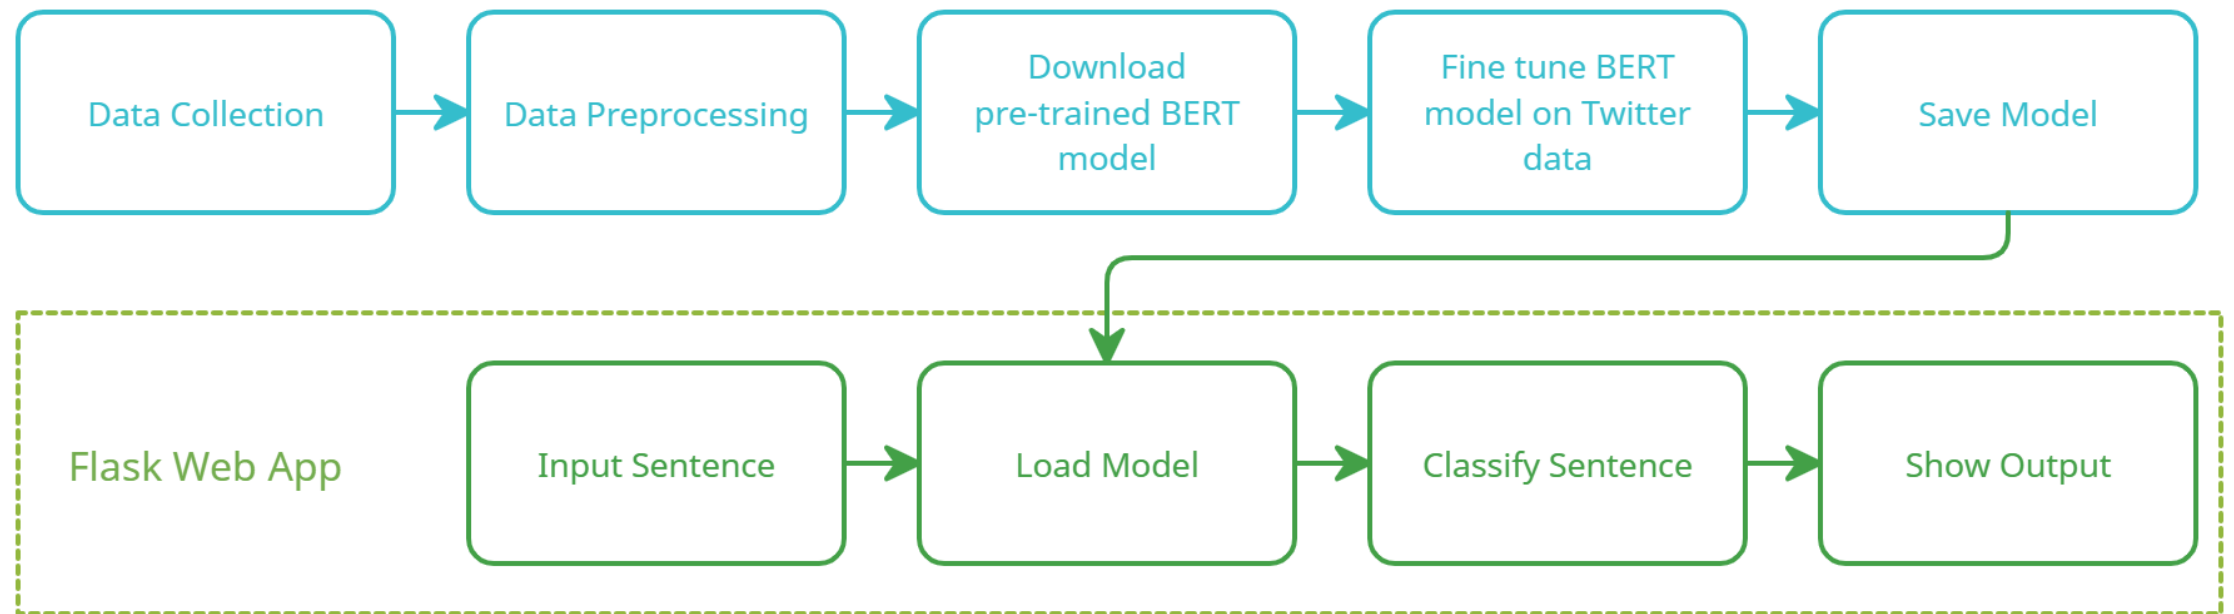
```
0.8894182002608317
```

F1 Score: Ensemble Model

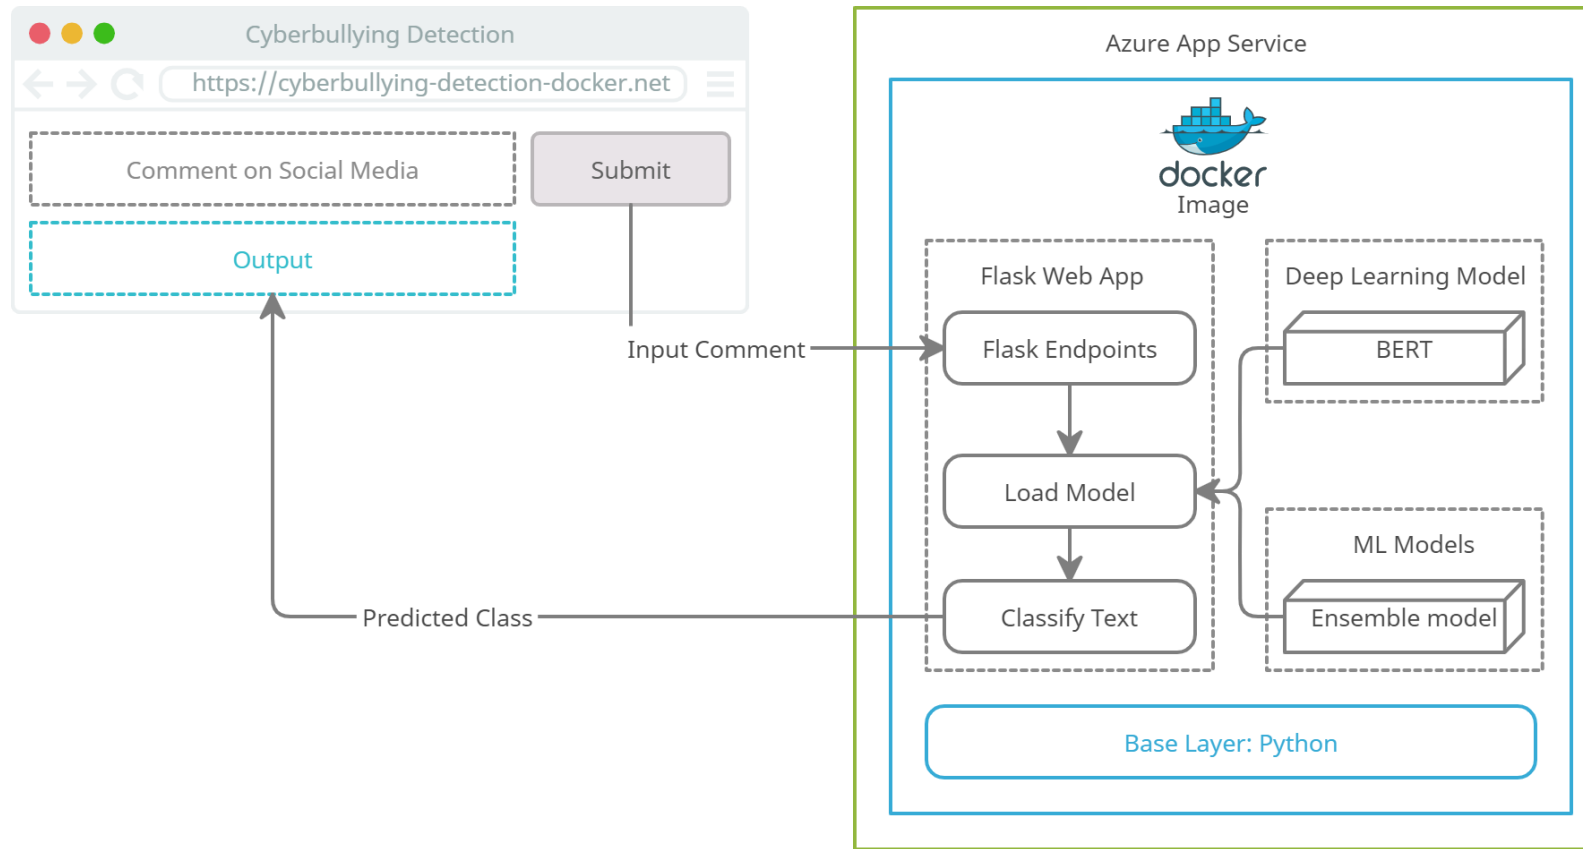
```
#Train and evaluation
for clf in (log_clf, svm_clf, mnb_clf, voting_clf):
    clf.fit( X_train_counts, y_train )
    y_pred = clf.predict(X_test_counts)
    print(clf.__class__.__name__, f1_score(y_test, y_pred, average = 'macro'))
```

```
LogisticRegression 0.7774494137250483
SVC 0.7791690204159132
MultinomialNB 0.7148849401810784
VotingClassifier 0.7848905890192499
```

Pipeline



Framework



Challenges

- Loading the BERT model failed with Flask Web App on Azure Linux
- I tried three different approaches to load the BERT model, but none of them worked on Azure Linux Machine. On the other hand, all three ways work on the local machine
- Flask App on Azure: [Cyberbullying Detection \(cyberbullying-detection.azurewebsites.net\)](https://cyberbullying-detection.azurewebsites.net/)

```
import torch
from transformers import BertForSequenceClassification, BertTokenizer, BertConfig

# -----Method 1-----#
# Save Model
model_to_save = model.module if hasattr(model, 'module') else model
model_to_save.save_pretrained(folder_path)
tokenizer.save_pretrained(folder_path)
# Load Model
model = BertForSequenceClassification.from_pretrained(folder_path)
tokenizer = BertTokenizer.from_pretrained(folder_path)

# -----Method 2-----#
# Save Model
torch.save(model, folder_path)
# Load Model
model = torch.load(folder_path)

# -----Method 3-----#
# Save Model
model_to_save = model.module if hasattr(model, 'module') else model
torch.save(model_to_save.state_dict(), "pytorch_model.bin")
model_to_save.config.to_json_file("config.json")
tokenizer.save_vocabulary("vocab.txt")
# Load Model
config = BertConfig.from_json_file("config.json")
model = BertForSequenceClassification(config)
state_dict = torch.load("pytorch_model.bin")
model.load_state_dict(state_dict)
tokenizer = BertTokenizer("vocab.txt", do_lower_case=True)
```

Solution

- Deployed the Flask App to Azure as a Docker container
- Docker image of Flask App on Azure:
[Cyberbullying Detection \(cyberbullying-detection-docker.azurewebsites.net\)](https://cyberbullying-detection-docker.azurewebsites.net)

Future Scope

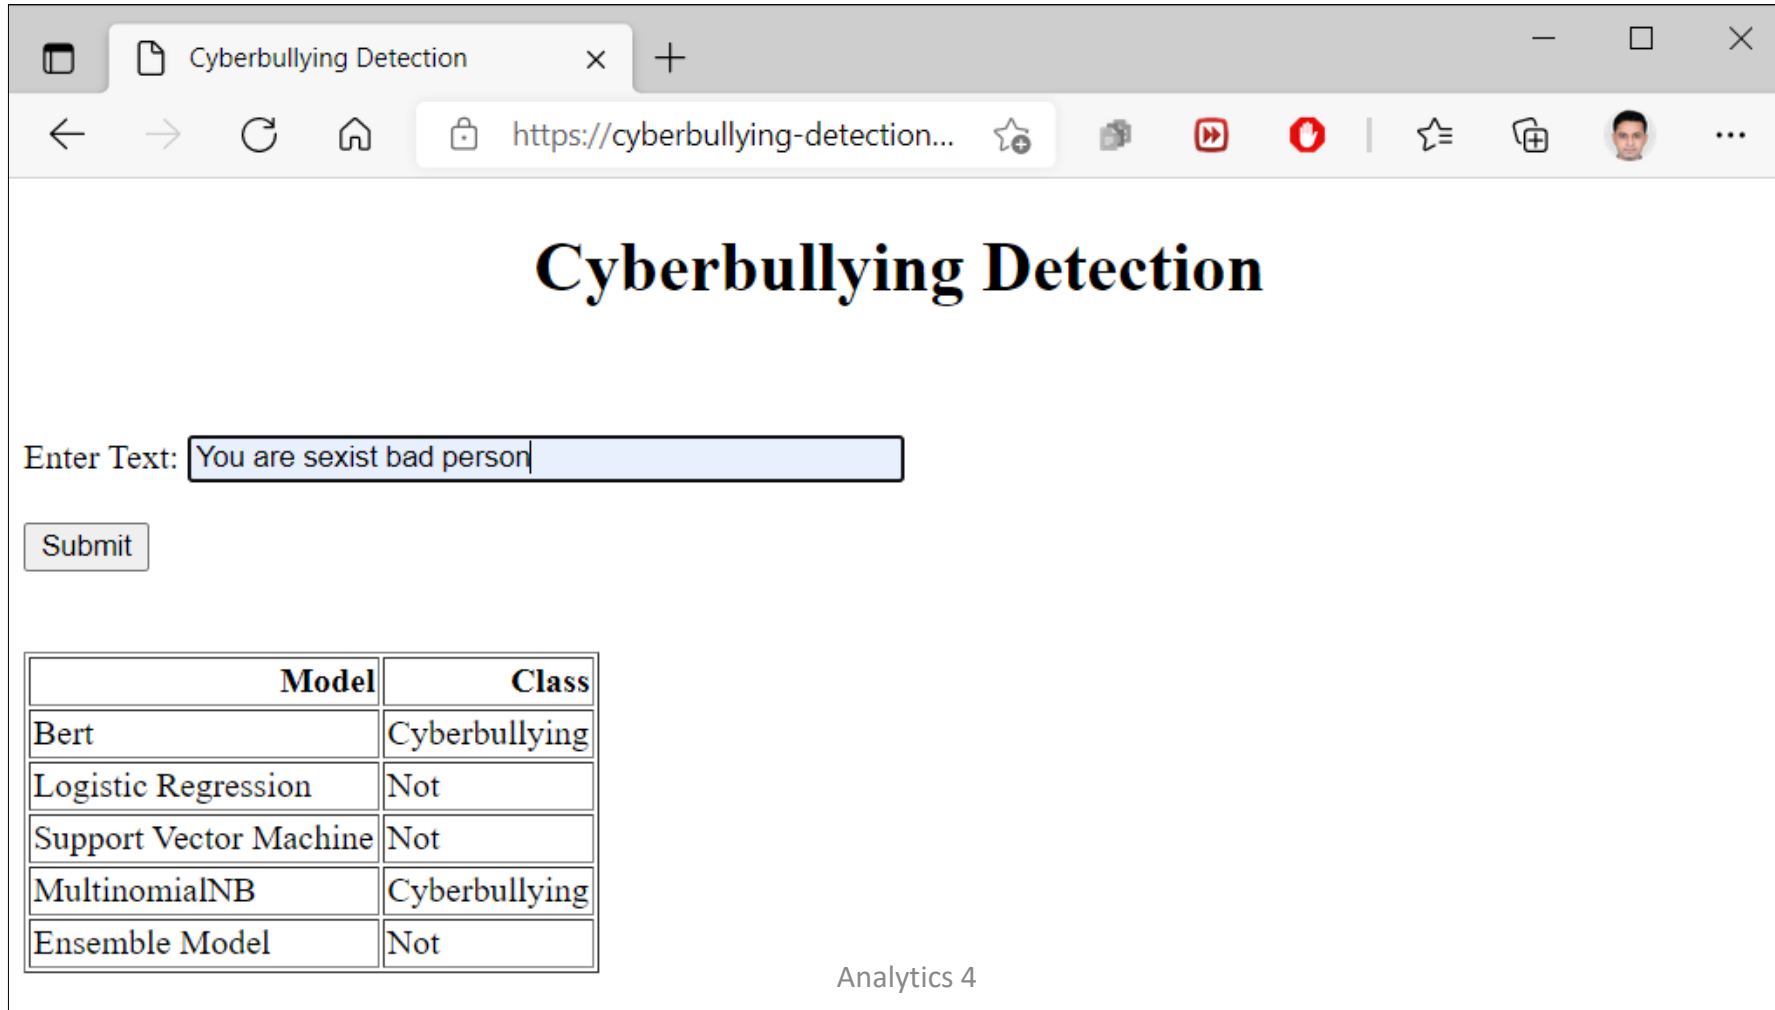
- Training BERT on data with more than two classes
- Comparison of BERT with Ensemble model for more than 2 classes

References

- Dataset 1: [HASOC \(hasocfire.github.io\)](https://hasocfire.github.io)
- Dataset 2: [UMICH SI650 - Sentiment Classification | Kaggle](#)
- [What Is Cyberbullying | StopBullying.gov](#)
- [What is BERT | BERT For Text Classification \(analyticsvidhya.com\)](#)
- [BERT Fine-Tuning Tutorial with PyTorch · Chris McCormick \(mccormickml.com\)](#)
- [BERT for Natural Language Processing | All You Need to know about BERT \(analyticsvidhya.com\)](#)

Demo

<https://cyberbullying-detection-docker.azurewebsites.net/>



The screenshot shows a web browser window with the title "Cyberbullying Detection". The address bar shows the URL "https://cyberbullying-detection...". The main heading is "Cyberbullying Detection". Below the heading, there is a text input field labeled "Enter Text:" containing the text "You are sexist bad person". A "Submit" button is located below the input field. At the bottom, there is a table showing the results of the detection for different models.

Model	Class
Bert	Cyberbullying
Logistic Regression	Not
Support Vector Machine	Not
MultinomialNB	Cyberbullying
Ensemble Model	Not

Thank You!